

Accountable Agents

An agent IDs initiative hosted by the [Center for AI Safety](#)

Framework Summary

Agents increasingly act on behalf of people and firms in transactions where neither the counterparty nor other impacted parties have any way to establish who, if anyone, is accountable for the actions of the agent.

The Accountable Agents framework is anchored on pairwise pseudonymous identifiers (PPIDs)—unique identifiers that allow an agent to be tied to a responsible legal person, or its principal, without revealing the identity of that person. Like a license plate, the identifier would not mean much on its own, but would allow counterparties and third parties to trace the agent back to a legal person as necessary.

A separate agent profile holds the substantive context a counterparty needs to decide whether to allow or disallow an agent's actions, such as authorization scope, an "agent bill of materials" (ABOM) covering the supply chain behind the agent, and context-specific fields such as intent. Profile data, principal data, and linkages to PPIDs may be preserved by record keepers, potentially incorporating data escrows, and disclosed in whole or in part in response to due process of law.

Verification is optional and tiered: a counterparty may ignore identity entirely, confirm only that a PPID exists, confirm that the PPID is tethered to a principal chain, and/or query specific pieces of the agent profile — at transaction time or later during an audit or incident investigation. The technical components that enable these functions are also largely extant, and may not require new wire protocol.

Working Plan

We aim to publish an accessible ~8-page arXiv paper detailing this framework to share with the broader standards and trust infrastructure community. The paper is protocol-neutral and technology-agnostic by design: it does not pick winners among protocols, and its goal is to align concepts and terminology between policymakers and technology implementers. Planned sections in the paper include:

1. **Introduction and problem:** introducing the identity challenge for agents, and the societal harms we seek to prevent
2. **Design requirements:** normative requirements (MUST/SHOULD/MAY) without picking specific protocols
3. **Framework overview:** framework overview, definitions, and existing work

4. **Worked scenarios:** in consequential settings, including financial services, healthcare, government services, and payments
5. **Conclusion and downstream considerations** — legal, governance, security, and adoption

Timeline

Milestone	Date
Agree on skeleton	Sep 18, 2026
Finalize author list	Oct 2, 2026
Finalize paper	Oct 16, 2026
Publish paper	Oct 31, 2026

People

Listed alphabetically. Core contributors have attended at least one meeting / contributed significant comments.

Core Contributors: Adam Khoja , Arunim Agarwal , cmturner@uga.edu , Dan Hendrycks , dazza@civics.com , ghadfield@jhu.edu , ian@credo.ai , idosalomon@gmail.com , Leo Wu , sam@aisafety.sg , wayne@spruceid.com

christian.schroeder@eng.ox.ac.uk , helen.toner@georgetown.edu , jfox@schmidtsciences.org , noam.kolt@mail.huji.ac.il , madhu@partnershiponai.org , Richard Ren

Get involved

This project came out of a working group at the Multi-Agent Ecosystems workshop hosted by the Center for AI Safety and Gillian Hadfield on August 8–9, 2026 at Stanford's CASBS ([shareout here](#)).

We are looking for both **co-authors** and **reviewers** — particularly from frontier labs, standards bodies, financial services, and policy.

For interest or questions, contact Leo Wu (leo@safe.ai).